

CS 343 Winter 2012 – Assignment 3

Instructor: Peter Buhr

Due Date: Monday, February 13, 2012 at 22:00

Late Date: Wednesday, February 15, 2012 at 22:00

February 12, 2012

This assignment introduces locks in $\mu\text{C++}$ and examines synchronization and mutual exclusion. Use it to become familiar with these new facilities, and ensure you use these concepts in your assignment solution. **(Tasks may have only public constructors and/or destructors; no other public members are allowed.)**

1. (a) Implement a generalized FIFO bounded-buffer for a producer/consumer problem with the following interface (you may add only a public destructor and private members):

```
template<typename T> class BoundedBuffer {
public:
    BoundedBuffer( const unsigned int size = 10 );
    void insert( T elem );
    T remove();
};
```

which creates a bounded buffer of size `size`, and supports multiple producers and consumers. You may *only* use `uCondLock` and `uOwnerLock` to implement the necessary synchronization and mutual exclusion needed by the bounded buffer.

Implement the `BoundedBuffer` in the following ways:

- i. Use busy waiting when waiting for buffer entries to become free or empty. In this approach, new tasks may barge into the buffer taking free or empty entries from tasks that have been signalled to access these entries. This implementation uses two condition-locks, one for each of the waiting producer and consumer tasks. The reason there is barging in this solution is that `uCondLock::wait` re-acquires its argument owner-lock before returning. Now once the owner-lock is released by a task exiting `insert` or `remove`, there is a race to acquire the lock by a new task calling `insert/remove` and by a signalled task. If the calling task wins the race, it barges ahead of any signalled task. So the state of the buffer at the time of the signal is not the same as the time the signalled task re-acquires the argument owner-lock, because the barging task changes the buffer. Hence, the signalled task may have to wait again, and there is no bound on how many times this may happen (possible long-term starvation).
- ii. Use no busy waiting when waiting for buffer entries to become free or empty. In this approach, new (barging) tasks must be prevented from taking free or empty entries if tasks have been unblocked to access these entries. This implementation uses three condition-locks, one for each of the waiting producer, consumer and barging tasks (*and has no looping*). **To prevent barging, create a flag variable to indicate when signalling is occurring; tasks entering the monitor check the flag to determine if they are barging because another task has been signalled.** Set the flag before signalling a producer/consumer task; ~~after the signalled task unblocks, it resets~~ the flag **is reset** only if there are no waiting **barging** tasks. If a task enters the buffer before a signal task has restarted (while the flag is set), it waits on the barging condition-lock. If a task needs to wait on either the producer/consumer condition-locks and the barging condition-lock is not empty, signal a task from it. Before exiting from the buffer and not signalling a producer/consumer, if the barging condition-lock is not empty, signal a task from it. Try to minimize the number of sets/resets of the signal flag.

Before inserting or removing an item to/from the buffer, perform an `assert` that checks if the buffer is not full or not empty, respectively. The buffer implementations are defined in a single file denoted in the following way:

```

#ifdef BUSY                                // busy waiting implementation
// implementation
#endif // BUSY

#ifdef NOBUSY                              // no busy waiting implementation
// implementation
#endif // NOBUSY

```

Test the bounded buffer with a number of producers and consumers. (Assume the buffer-element type, `T`, is `int` and the sentinel value is -1.) The producer interface is:

```

_Task Producer {
    void main();
public:
    Producer( BoundedBuffer<int> &buffer, const int Produce, const int Delay );
};

```

The producer generates `Produce` integers, from 1 to `Produce` inclusive, and inserts them into `buffer`. Before producing an item, a producer randomly yields between 0 and `Delay-1` times. In μ C++, yielding is accomplished by calling `yield(unsigned int)`. The consumer interface is:

```

_Task Consumer {
    void main();
public:
    Consumer( BoundedBuffer<int> &buffer, const int Delay, const int Sentinel, int &sum );
};

```

The consumer removes items from `buffer`, and terminates when it removes a `Sentinel` value from the buffer. A consumer sums all the values it removes from `buffer` (excluding the `Sentinel` value) and returns this value through the reference variable `sum`. Before removing an item, a consumer randomly yields between 0 and `Delay-1` times.

`uMain::main` creates the bounded buffer, the producer and consumer tasks. After all the producer tasks have terminated, `uMain::main` inserts an appropriate number of sentinel values into the buffer to terminate the consumers. The partial sums from each consumer are totalled to produce the sum of all values generated by the producers. Print this total in the following way:

```
total: ddddd
```

The sum must be the same regardless of the order or speed of execution of the producer and consumer tasks.

The shell interface for the `boundedBuffer` program is:

```
boundedBuffer [ Cons [ Prods [ Produce [ BufferSize [ Delays ] ] ] ] ]
```

(Square brackets indicate optional command line parameters, and do not appear on the actual command line.) Where the meaning of each parameter is:

Cons: positive number of consumers to create. The default value if unspecified is 5.

Prods: positive number of producers to create. The default value if unspecified is 3.

Produce: positive number of items generated by each producer. The default value if unspecified is 10.

BufferSize: positive number of elements in (size of) the bounded buffer. The default value if unspecified is 10.

Delays: positive number of times a producer/consumer yields *before* inserting/removing an item into/from the buffer. The default value if unspecified is `Cons + Prods`.

Check all command arguments for correct form (integers) and range; print an appropriate usage message and terminate the program if a value is missing or invalid. Use the C library routine `rand_r` to safely generate random values. Each producer and consumer task creates their own seed, which is initialized to a value so each run of the program generates different orderings of delays.

- (b) i. Compare the busy and non-busy waiting versions of the program with respect to performance by doing the following:

- Time the execution using the time command:

```
% time ./a.out
3.21u 0.02s 0:03.32 100.0%
```

(Output from time differs depending on your shell, but all provide user, system and real time.) Compare the *user* time (3.21u) only, which is the CPU time consumed solely by the execution of user code (versus system and real time).

- Use the program command-line arguments 50 55 10000 30 10 and adjust the Produce amount (if necessary) to get execution times in the range 0.1 to 100 seconds. (Timing results below 0.1 seconds are inaccurate.) Use the same command-line values for all experiments.
 - Run both the experiments again after recompiling the programs with compiler optimization turned on (i.e., compiler flag -O2). Include all 4 timing results to validate your experiments.
- State the observed performance difference between busy and nobusy waiting execution, without and with optimization.
 - Speculate as to the reason for the performance difference between busy and nobusy waiting execution.
 - Add the following declaration to uMain::main after checking command-line arguments but before creating any tasks:

```
#ifdef __U_MULTI__
    uProcessor p[3] __attribute__(( unused )); // create 3 kernel thread for a total of 4
#endif // __U_MULTI__
```

to increase the number of kernel threads to access multiple processors. This declaration must be in the same scope as the declaration of the quicksort task. Compile the program with the -multi flag and no optimization on a multi-core computer with at least 4 CPUs (cores), and run the same experiment as above. Include timing results to validate your experiment.

- State the observed performance difference between uniprocessor and multiprocessor execution.
 - Speculate as to the reason for the performance difference between uniprocessor and multiprocessor execution.
- Consider the following situation involving a tour group of V tourists. The tourists arrive at the Louvre Museum for a tour. However, a tour can only be composed of G people at a time, otherwise the tourists cannot hear what the guide is saying. As well, there are 2 kinds of tours available at the Louvre: picture and statue art. Therefore, each group of G tourists must vote amongst themselves to select the kind of tour they want to take. During voting, tasks must block until all votes are cast, i.e., assume a secret ballot. Once a decision is made, the tourists in that group proceed on the tour.

To simplify the problem, the program only has to handle cases where the number of tourists in a group is odd, precluding a tie vote, and the total number of tourists is evenly divisible by the tour-group size so all groups contain the same number of tourists.

Implement a vote-tallier for G -way voting as a class using *only* $\mu\text{C++}$:

- uCondLock and uOwnerLock to provide mutual exclusion and synchronization,
- uSemaphore to provide mutual exclusion and synchronization.

No busy waiting is allowed and barging tasks can spoil an election and must be prevented. Figure 1 shows the different forms for each $\mu\text{C++}$ vote-tallier implementation (you may add only a public destructor and private members), where the preprocessor is used to conditionally compile a specific interface. This form of header file removes duplicate code. When the vote-tallier is created, it is passed the size of a group and a printer for printing state transitions. There is only one vote-tallying object created for all of the voters, who share a reference to it. Each tourist task calls the vote method with their id and a vote of either true or false, indicating their desire for the picture or statue tour, respectively. The vote routine does not return until group votes are cast; after which, the majority result of the voting (true or false) is returned to each voter. The groups are formed based on voter arrival; e.g., for a group of 3, if voters 2, 5, 8 cast their votes first, they form the first group, etc. Hence, all voting is serialized. An appropriate preprocessor variable is defined on the compilation command using the following syntax:

```
u++ -DIMPLTYPE_SEM -c TallyVotesSEM.cc
```

```

#if defined( IMPLTYPE_MC )           // mutex/condition solution
// includes for this kind of vote-tallier
class TallyVotes {
    // private declarations for this kind of vote-tallier
#elif defined( IMPLTYPE_SEM )       // semaphore solution
// includes for this kind of vote-tallier
class TallyVotes {
    // private declarations for this kind of vote-tallier
#else
    #error unsupported voter type
#endif
    // common declarations
public:                             // common interface
    TallyVotes( unsigned int group, Printer &printer );
    bool vote( unsigned int id, bool ballot );
};

```

Figure 1: Tally Voter Interfaces

The interface for the voting task is (you may add only a public destructor and private members):

```

_Task Voter {
public:
    enum States { Start = 'S', Vote = 'V', Block = 'B', Unblock = 'U', Complete = 'C', Finished = 'F' };
    Voter( unsigned int id, TallyVotes &voteTallier, Printer &printer );
};

```

The task main of a voting task looks like:

- print start message
- yield a random number of times, between 0 and 19 inclusive, so all tasks do not start simultaneously
- randomly calculate a ballot (true or false)
- vote at the vote-tallier

Note that each task votes only once. In μ C++, yielding is accomplished by calling `yield(unsigned int)`.

All output from the program is generated by calls to a printer, excluding error messages. The interface for the printer is (you may add only a public destructor and private members) (monitors will be discussed shortly):

```

_Monitor / _Cormonitor Printer {    // choose one of the two kinds of type constructor
public:
    Printer( unsigned int voters );
    void print( unsigned int id, Voter::States state );
    void print( unsigned int id, Voter::States state, bool vote );
    void print( unsigned int id, Voter::States state, unsigned int numBlocked );
};

```

The printer attempts to reduce output by storing information for each voter until one of the stored elements is overwritten. When information is going to be overwritten, all the stored information is flushed and storing starts again. Output must look like that in Figure 2.

Each column is assigned to a voter with an appropriate title, e.g., “Voter0”, and a column entry indicates its current status:

State	Meaning
S	starting
V <i>b</i>	voting with ballot <i>b</i> (0/1)
B <i>n</i>	blocking during voting, <i>n</i> voters waiting (including self)
U <i>n</i>	unblocking after group reached, <i>n</i> voters still waiting (not including self)
C	group is complete and voting result is computed
F <i>b</i>	finished voting and result of vote is <i>b</i> (0/1)

Information is buffered until a column is overwritten for a particular entry, which causes the buffered data to be flushed. If there is no new stored information for a column since the last buffer flush, an empty column is

1	% vote 3 1		
2	Voter0	Voter1	Voter2
3	=====		
4	S	S	S
5			V 0
6			C
7	F 0
8	V 1		
9	C		
10	F 1	...	
11		V 1	
12		C	
13	...	F 1	
14	=====		
15	All tours started		

1	% vote 6 3					
2	Voter0	Voter1	Voter2	Voter3	Voter4	Voter5
3	=====					
4	S	S				
5		V 1				
6		B 1	S	S	S	S
7			V 1			
8	V 1		B 2			
9	C					
10	F 1
11			U 1			
12	F 1
13		U 0				
14	...	F 1
15						V 0
16				V 0		B 1
17				B 2	V 1	
18					C	
19	F 0	...
20				U 1		
21	F 0
22						U 0
23	F 0
24	=====					
25	All tours started					

Figure 2: Voters: Example Output

printed. When a task finishes, the buffer is flushed immediately, the state for that object is marked with F, and all other objects are marked with "...". After a task has finished, no further output appears in that column. All output spacing can be accomplished using the standard 8-space tabbing. Buffer any information necessary for printing in internal representation; **do not build and store strings of text for output**. Calls to perform printing may be performed from the vote-tallier and/or a voter task (you decide where to print).

For example, in line 4 of the right hand example in Figure 2, Voter1 has the value "S" in its buffer slot, and the first buffer slot is full and the last 4 slots are empty. When Vote1 attempts to print "V 1", which overwrites its current buffer value of "S", the buffer must be flushed generating line 4. Voter1's new value of "V 1" is then inserted into its buffer slot. When Vote1 attempts to print "B 1", which overwrites its current buffer value of "V 1", the buffer must be flushed generating line 5 and no other values are printed on the line because the print is consecutive (i.e., no intervening call from another object).

The executable program is named vote and has the following shell interface:

```
vote [ V [ G [ Seed ] ] ]
```

V is the size of a tour, i.e., the number of voters (tasks) to be started (multiple of G); if V is not present, assume a value of 6. G is the size of a tour group (odd number); if G is not present, assume a value of 3. Seed is the seed for the random-number generator and must be greater than 0. If the seed is unspecified, use a random value like the process identifier (getpid) or current time (time), so each run of the program generates different output. Use the following monitor to safely generate random values (monitors will be discussed shortly):

```
_Monitor PRNG {
public:
    PRNG( unsigned int seed = 1009 ) { srand( seed ); } // set seed
    void seed( unsigned int seed ) { srand( seed ); } // set seed
    unsigned int operator()() { return rand(); } // [0,UINT_MAX]
    unsigned int operator()( unsigned int u ) { return operator()() % (u + 1); } // [0,u]
    unsigned int operator()( unsigned int l, unsigned int u ) { return operator()( u - l ) + l; } // [l,u]
};
```

Note, because of the non-deterministic execution of concurrent programs, multiple runs with a common seed may not generate the same output. Nevertheless, shorts runs are often the same so the seed can be useful for

testing. Check all command arguments for correct form (integers) and range; print an appropriate usage message and terminate the program if a value is missing or invalid.

Submission Guidelines

Please follow these guidelines carefully. Review the [Assignment Guidelines](#) and [C++ Coding Guidelines](#) *before* starting each assignment. **Each text file, i.e., *.txt file, must be ASCII text and not exceed 500 lines in length, where a line is 120 characters.** Programs should be divided into separate compilation units, i.e., *.{h,cc,C,cpp} files, where applicable. Use the [submit](#) command to electronically copy the following files to the course account.

1. q1buffer.h, q1*.{h,cc,C,cpp} – code for question question [1a, p. 1](#). **Program documentation must be present in your submitted code. No test, user or system documentation is to be submitted for this question. Output for this question is checked via a marking program, so it must match exactly with the given program.**
2. q1buffer.txt – contains the information required by question [1b, p. 2](#).
3. q2tallyVotes.h, q2*.{h,cc,C,cpp} – code for question question [2, p. 3](#). **Program documentation must be present in your submitted code. No user or system documentation is to be submitted for this question.**
4. q2vote.testtxt – test documentation for question [2, p. 3](#), which includes the input and output of your tests, documented by the use of script and after being formatted by scriptfix. **Poor documentation of how and/or what is tested can result in a loss of all marks allocated to testing.**
5. Use the following Makefile to compile the programs for question [1a, p. 1](#) and [2, p. 3](#):

```
KIND:=NOBUSY
OPT:=
TYPE:=SEM

CXX = u++                                # compiler
CXXFLAGS = -g -Wall -Wno-unused-label -MMD ${OPT} -D"${KIND}" -D"IMPLTYPE_${TYPE}"
MAKEFILE_NAME = ${firstword ${MAKEFILE_LIST}} # makefile name

OBJECTS1 = # list of object files for question 1 prefixed with "q1"
EXEC1 = boundedBuffer

OBJECTS2 = q2tallyVotes${TYPE}.o # list of object files for question 2 prefixed with "q2"
EXEC2 = vote

OBJECTS = ${OBJECTS1} ${OBJECTS2}        # all object files
DEPENDS = ${OBJECTS:.o=.d}               # substitute ".o" with ".d"
EXECS = ${EXEC1} ${EXEC2}                # all executables

.PHONY : all clean

all : ${EXECS}                            # build all executables

#####

-include ImplKind

ifeq (${IMPLKIND},${KIND})                # same implementation type as last time ?
${EXEC1} : ${OBJECTS1}
    ${CXX} ${CXXFLAGS} $^ -o $@
else
ifeq (${KIND},)                          # no implementation type specified ?
# set type to previous type
KIND=${IMPLKIND}
${EXEC1} : ${OBJECTS1}
    ${CXX} ${CXXFLAGS} $^ -o $@
else
# implementation type has changed
.PHONY : ${EXEC1}
```

```

${EXEC1} :
    rm -f ImplKind
    touch q1buffer.h
    ${MAKE} ${EXEC1} KIND=${KIND}
endif
endif

ImplKind :
    echo "IMPLKIND=${KIND}" > ImplKind

#####

-include ImplType

ifeq (${IMPLTYPE},${TYPE})                                # same implementation type as last time ?
${EXEC2} : ${OBJECTS2}
    ${CXX} ${CXXFLAGS} $^ -o $@
else
ifeq (${TYPE},)                                            # no implementation type specified ?
    # set type to previous type
    TYPE=${IMPLTYPE}
    ${EXEC2} : ${OBJECTS2}
        ${CXX} ${CXXFLAGS} $^ -o $@
else                                                        # implementation type has changed
.PHONY : ${EXEC2}
${EXEC2} :
    rm -f ImplType
    touch q2tallyVotes.h
    ${MAKE} ${EXEC2} TYPE="${TYPE}"
endif
endif

ImplType :
    echo "IMPLTYPE=${TYPE}" > ImplType

#####

${OBJECTS} : ${MAKEFILE_NAME}                            # OPTIONAL : changes to this file => recompile

-include ${DEPENDS}

clean :                                                    # remove files that can be regenerated
    rm -f *.d *.o ${EXECS} ImplKind ImplType

This makefile is used as follows:

$ make boundedBuffer KIND=BUSY
$ boundedBuffer ...
$ make boundedBuffer KIND=NOBUSY OPT='-multi -O2'
$ boundedBuffer ...
$
$ make vote TYPE=MC
$ vote ...
$ make vote TYPE=SEM
$ vote ...

```

Put this Makefile in the directory with the programs, name the source files as specified above, and then type make buffer or make vote in the directory to compile the programs. This Makefile must be submitted with the assignment to build the program, so it must be correct. Use the web tool [Request Test Compilation](#) to ensure you have submitted the appropriate files, your makefile is correct, and your code compiles in the testing environment. **If the makefile fails or does not produce correctly named executables, or if a program does not compile, you receive zero for all “Testing” marks.**

Follow these guidelines. Your grade depends on it!